

Gaurav Chaudhari

✉ chaudharigaurav37@gmail.com | 📞 +91 8855873204 | 📍 Mumbai, India

🌐 [linkedin.com/in/gaurav-chaudhari-gc](https://www.linkedin.com/in/gaurav-chaudhari-gc) | 🐙 github.com/gaurav-163

Experience

Anvex AI

Sept 2025 – Present

Jr. Software Engineer ML

Vashi, India

- Architected high-throughput inference microservices using **FastAPI** on **Docker**, reducing model latency to **<200ms** for real-time interaction.
- Engineered **Form2Dial**, a scalable voice agent pipeline integrating **VAPI SDK** and **SIP Trunking** to manage concurrent telephony streams with dynamic context injection.
- Implemented strict JSON schema enforcement for LLM outputs, increasing structured data extraction accuracy by **75%** and eliminating hallucinated fields.
- Built **Anvex Voice**, a production-grade **RAG system** that retrieves document-grounded context in real-time, reducing agent response time during live calls.
- Fine-tuned SLMs and **NVIDIA Parakeet TTS** models on proprietary datasets, optimizing inference cost and reducing Word Error Rate (WER) for domain vocabulary.

Anvex AI

Jul 2025 – Sept 2025

AI/ML Engineer Intern

Vashi, India

- Designed and evaluated **20+ Chain-of-Thought (CoT)** prompt templates, ensuring deterministic behavior for enterprise-facing AI agents.
- Optimized conversational workflows using user-level analytics, achieving a **30% uplift in engagement** and a **15% reduction in call drop-offs**.

Technical Skills

Generative AI: LLM Fine-Tuning (Unsloth, PEFT/LoRA), RAG, LangChain, CrewAI, Prompt Engineering (CoT), OpenAI API

Audio & Multimodal: VAPI, Deepgram, Sarvam AI, TTS, ASR

Machine Learning: PyTorch, TensorFlow, Scikit-Learn, Transformers, Vector Search (FAISS, ChromaDB).

Infrastructure & DevOps: Python, FastAPI, Docker, Kafka, GCP, GitHub Actions.

Data: MySQL, MongoDB, Pandas, NumPy.

Projects

CavScan: Deep Learning Diagnostic System | ViT, Computer Vision

Jun 2024 – Mar 2025

- Developed an end-to-end dental cavity detection pipeline using **Vision Transformers (ViT)** trained on 20K X-rays, achieving a **0.91 F1-score**.
- Optimized inference performance for edge deployment, achieving **<200ms latency**, and added Grad-CAM for explainability.

Mira: Knowledge-Based RAG Assistant | ChromaDB, LangChain, Hybrid Search

Jan 2025 – Feb 2025

- Engineered a domain-agnostic RAG pipeline utilizing **Hybrid Search (Dense + Keyword)** to retrieve context from unstructured PDFs and handwritten notes via OCR.
- Integrated **Cross-Encoder Reranking** to filter retrieved chunks, reducing hallucination rates by ensuring high context relevance before LLM generation.
- Deployed the retrieval microservice using FastAPI, optimizing vector search latency to **<100ms** on large document stores.

News Recommender System | NLP, Flask, REST API

Apr 2023 – Oct 2023

- Built a News Recommender System using **Sentence Transformers** for semantic embeddings and **FAISS** for efficient similarity search, deployed with a **Flask**-based interface for real-time personalized article retrieval.

Education

New Horizon Institute Of Technology & Management

Mumbai, India

B.E. in Artificial Intelligence and Data Science

2021 – 2025

- Global Winner:** LabLab.ai Fine-Tuning Hackathon — 1st Place (LLM optimization).
- Leadership:** Head Technical Secretary, Student Association.

Certifications

DeepLearning.AI: Generative AI with Large Language Models | **Stanford/Coursera:** Machine Learning Specialization